

Perform Foundational Data, ML, and AI Tasks in Google Cloud Challenge Lab

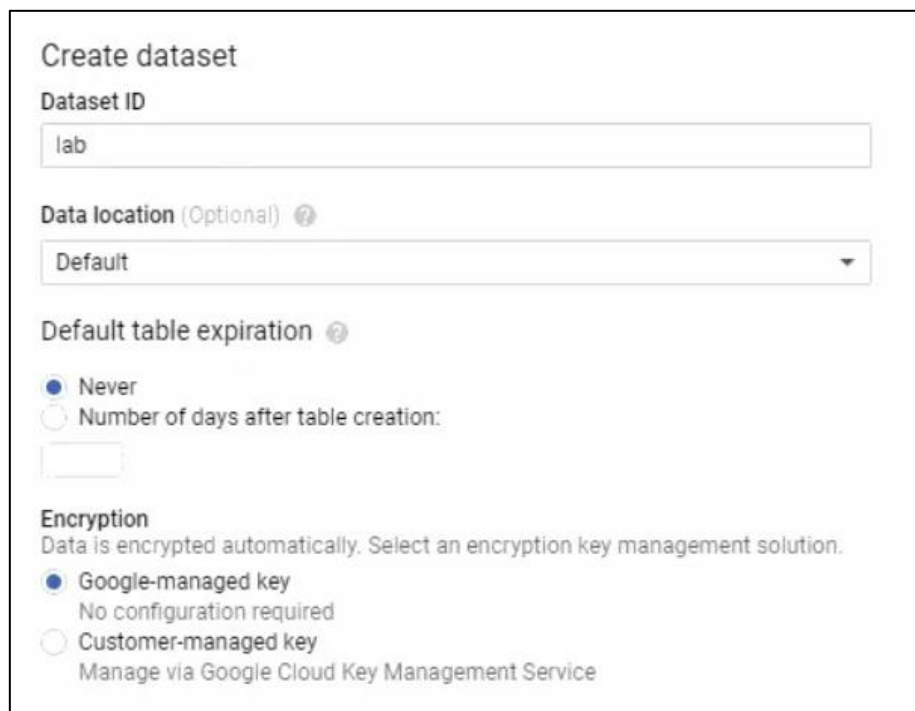
Author: Vedant Kakde | **GitHub Profile:** github.com/vedant-kakde | **LinkedIn Profile:** linkedin.com/in/vedant-kakde/

Task 1: Run a simple Dataflow job

In this task, you have to transfer the data in a CSV file to BigQuery using Dataflow via Pub/Sub. First of all, you need to create a BigQuery dataset called `lab` and a Cloud Storage bucket called with your project ID.

1.1 Created a BigQuery dataset called `Lab`

1. In the Cloud Console, click on **Navigation Menu > BigQuery**.
2. Select your project in the left pane.
3. Click **CREATE DATASET**.
4. Enter `lab` in the Dataset ID, then click **Create dataset**.



Create dataset

Dataset ID

lab

Data location (Optional) ?

Default

Default table expiration ?

☒ Never

☐ Number of days after table creation:

Encryption

Data is encrypted automatically. Select an encryption key management solution.

☒ Google-managed key

No configuration required

☐ Customer-managed key

Manage via Google Cloud Key Management Service

5. Run `gsutil cp gs://cloud-training/gsp323/lab.schema` in the Cloud Shell to download the schema file.
6. View the schema by running `cat lab.schema`.

```
[
  {"type": "STRING", "name": "guid"},
  {"type": "BOOLEAN", "name": "isActive"},
  {"type": "STRING", "name": "firstname"},
  {"type": "STRING", "name": "surname"},
  {"type": "STRING", "name": "company"},
  {"type": "STRING", "name": "email"},
  {"type": "STRING", "name": "phone"},
  {"type": "STRING", "name": "address"},
  {"type": "STRING", "name": "about"},
  {"type": "TIMESTAMP", "name": "registered"},
  {"type": "FLOAT", "name": "latitude"},
  {"type": "FLOAT", "name": "longitude"}
]
```

Author: Vedant Kakde | **GitHub Profile:** github.com/vedant-kakde | **LinkedIn Profile:** linkedin.com/in/vedant-kakde/

Copy this schema into notepad. We will use it in further step.

7. Go back to the Cloud Console, select the new dataset **lab** and click **Create Table**.
8. In the Create table dialog, select **Google Cloud Storage** from the dropdown in the Source section.
9. Copy `gs://cloud-training/gsp323/lab.csv` to **Select file from GCS bucket**.
10. Enter `customers` to "Table name" in the Destination section.
11. Enable **Edit as text** and copy the JSON data from the `lab.schema` file to the textarea in the Schema section.
12. Click **Create table**.

Create table

Source

Create table from: Google Cloud Storage Select file from GCS bucket: gs://cloud-training/gsp323/lab.csv Browse File format: CSV

☐ Source Data Partitioning

Destination

☒ Search for a project ☐ Enter a project name

Project name: qwiklabs-gcp-00-e66de107d4d4 Dataset name: lab Table type: Native table

Table name: customers

Schema

Auto detect ☒ Schema and input parameters

☒ Edit as text

```
1 [
2   {"type": "STRING", "name": "guid"},
3   {"type": "BOOLEAN", "name": "isActive"},
4   {"type": "STRING", "name": "firstName"},
5   {"type": "STRING", "name": "surname"},
6   {"type": "STRING", "name": "company"},
7   {"type": "STRING", "name": "email"},
8   {"type": "STRING", "name": "phone"},
9   {"type": "STRING", "name": "address"},
10  {"type": "STRING", "name": "about"},
11  {"type": "TIMESTAMP", "name": "registered"},
12  {"type": "FLOAT", "name": "latitude"},
13  {"type": "FLOAT", "name": "longitude"}
14 ]
```

Partition and cluster settings

Partitioning: No partitioning

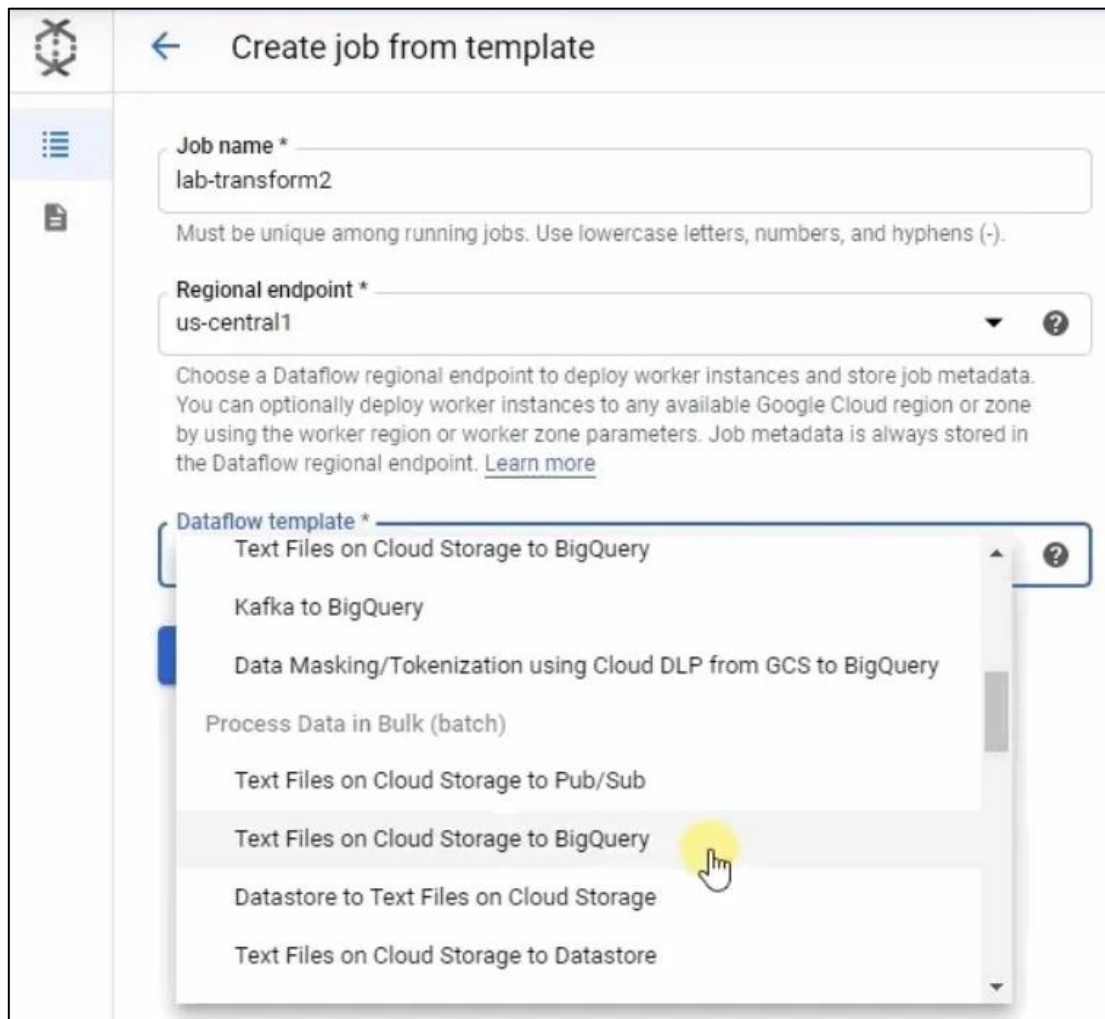
1.2 Create a Cloud Storage bucket

1. In the Cloud Console, click on **Navigation Menu > Storage**.
2. Click **CREATE BUCKET**.
3. Copy your GCP Project ID to Name your bucket.

4. Click **CREATE**.

1.3 Create a Dataflow job

1. In the Cloud Console, click on **Navigation Menu > Dataflow**.
2. Click **CREATE JOB FROM TEMPLATE**.
3. In Create job from template, give an arbitrary job name.
4. From the dropdown under Dataflow template, select **Text Files on Cloud Storage Pub/Sub** under “Process Data in Bulk (batch)”. (**DO NOT** select the item under “Process Data Continuously (stream)”).



Job name *

lab-transform2

Must be unique among running jobs. Use lowercase letters, numbers, and hyphens (-).

Regional endpoint *

us-central1

Choose a Dataflow regional endpoint to deploy worker instances and store job metadata. You can optionally deploy worker instances to any available Google Cloud region or zone by using the worker region or worker zone parameters. Job metadata is always stored in the Dataflow regional endpoint. [Learn more](#)

Dataflow template *

Text Files on Cloud Storage to BigQuery

Kafka to BigQuery

Data Masking/Tokenization using Cloud DLP from GCS to BigQuery

Process Data in Bulk (batch)

Text Files on Cloud Storage to Pub/Sub

Text Files on Cloud Storage to BigQuery

Datastore to Text Files on Cloud Storage

Text Files on Cloud Storage to Datastore

5. Under the Required parameters, enter the following values:

Field	Value
JavaScript UDF path in Cloud Storage	gs://cloud-training/gsp323/lab.js
JSON path	gs://cloud-training/gsp323/lab.schema
JavaScript UDF name	transform
BigQuery output table	YOUR_PROJECT:lab.customers

Field	Value
Cloud Storage input path	gs://cloud-training/gsp323/lab.csv
Temporary BigQuery directory	gs://YOUR_PROJECT/bigquery_temp
Temporary location	gs://YOUR_PROJECT/temp

Replace **YOUR_PROJECT** with your project ID.

Create job from template

Dataflow template *
Text Files on Cloud Storage to BigQuery

Batch pipeline. Reads text files stored in Cloud Storage, transforms them using a JavaScript user-defined function (UDF), and outputs the result to BigQuery.

Required parameters

JavaScript UDF path in Cloud Storage *
gs://cloud-training/gsp323/lab.js
The Cloud Storage path pattern for the JavaScript code containing your user-defined functions. Ex: gs://<your-bucket>/<your-transforms>/*.js

JSON path *
gs://cloud-training/gsp323/lab.schema
The Cloud Storage path to the JSON file that defines your BigQuery schema. Ex: gs://<your-bucket>/<your-schema>.json

JavaScript UDF name *
transform
The name of the function to call from your JavaScript file. Use only letters, digits, and underscores. Ex: transform_udf1

BigQuery output table *
qwiklabs-gcp-00-e66de107d4d4:lab.customers
The location of the BigQuery table in which to store your processed data. If you reuse an existing table, it will be overwritten. Ex: <your-project>:<your-dataset>:<your-table>

Cloud Storage input path *
gs://cloud-training/gsp323/lab.csv
The path to the Cloud Storage text to read. Ex: gs://<your-bucket>/<your-file>.txt

Temporary BigQuery directory *
gs://qwiklabs-gcp-00-e66de107d4d4/bigquery_temp
Temporary directory for the BigQuery loading process. Ex: gs://<your-bucket>/<your-files>/temp-dir

Temporary location *
gs://qwiklabs-gcp-00-e66de107d4d4/temp
Path and filename prefix for writing temporary files. Ex: gs://<your-bucket>/temp

Job components:

- JavascriptT...Javascript
- BigQueryCon... ToTableRow
- Insert into Bigquery
- DropInputs

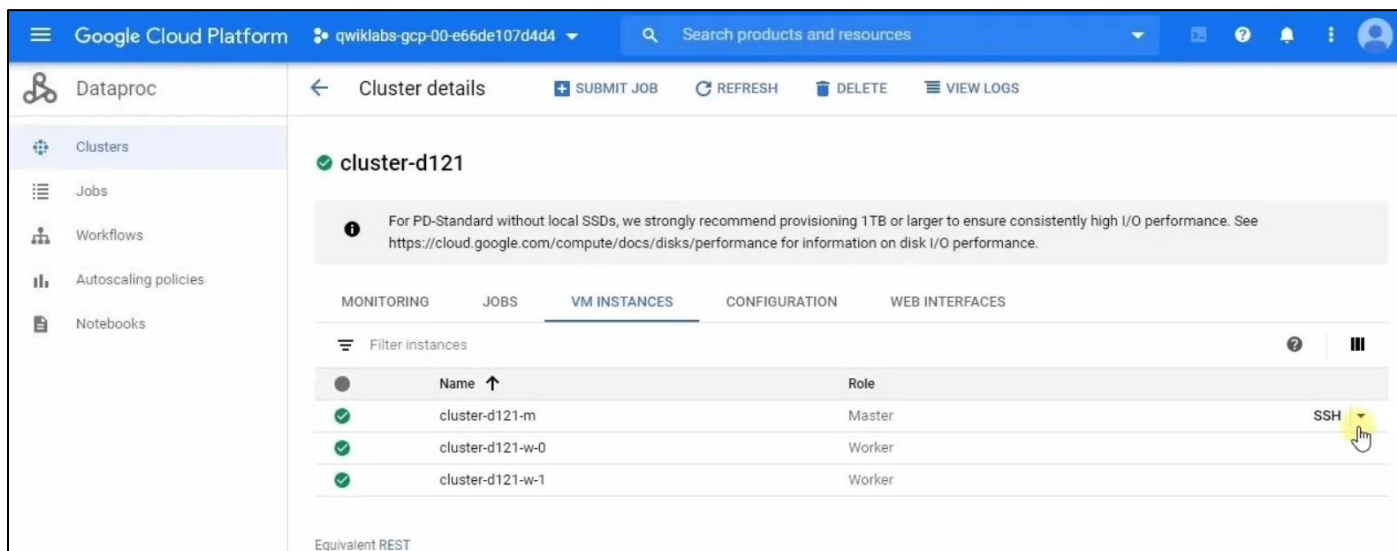
6. Click **RUN JOB**.

Task 2: Run a simple Dataproc job

Create a Dataproc cluster

1. In the Cloud Console, click on **Navigation Menu > Dataproc > Clusters**.
2. Click **CREATE CLUSTER**.
3. Make sure the cluster is going to create in the region **us-central1**.
4. Click **Create**.

5. After the cluster has been created, click the **SSH** button in the row of the master instance.



Google Cloud Platform | qwiklabs-gcp-00-e66de107d4d4 | Search products and resources

Dataproc | Cluster details | SUBMIT JOB | REFRESH | DELETE | VIEW LOGS

cluster-d121

For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure consistently high I/O performance. See <https://cloud.google.com/compute/docs/disks/performance> for information on disk I/O performance.

MONITORING | JOBS | VM INSTANCES | CONFIGURATION | WEB INTERFACES

Filter instances

Name ↑	Role
cluster-d121-m	Master
cluster-d121-w-0	Worker
cluster-d121-w-1	Worker

SSH

Equivalent REST

6. In the SSH console, run the following command:

```
hdfs dfs -cp gs://cloud-training/gsp323/data.txt /data.txt
```

7. Close the SSH window and go back to the Cloud Console.
8. Click **SUBMIT JOB** on the cluster details page.
9. Select **Spark** from the dropdown of “Job type”.
10. Copy `org.apache.spark.examples.SparkPageRank` to “Main class or jar”.
11. Copy `file:///usr/lib/spark/examples/jars/spark-examples.jar` to “Jar files”.
12. Enter `/data.txt` to “Arguments”.

Submit a job

Job type *
Spark

Main class or jar *
org.apache.spark.examples.SparkPageRank
The fully qualified name of a class in a provided or standard jar file, for example, com.example.wordcount, or a provided jar file to use the main class of that jar file

Jar files
file:///usr/lib/spark/examples/jars/spark-examples.jar
Enter file path, for example, hdfs://example/example.jar
Jar files are included in the CLASSPATH. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix.

Archive files
Archive files are extracted in the Spark working directory. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix. Supported file types: jar, tar, tar.gz, tgz, .zip.

Arguments
/data.txt
Additional arguments to pass to the main class. Press <Return> after each argument.

Max restarts per hour
Leave blank if you don't want to allow automatic restarts on job failure. [Learn more](#)

Properties ⓘ
[+ ADD PROPERTY](#)

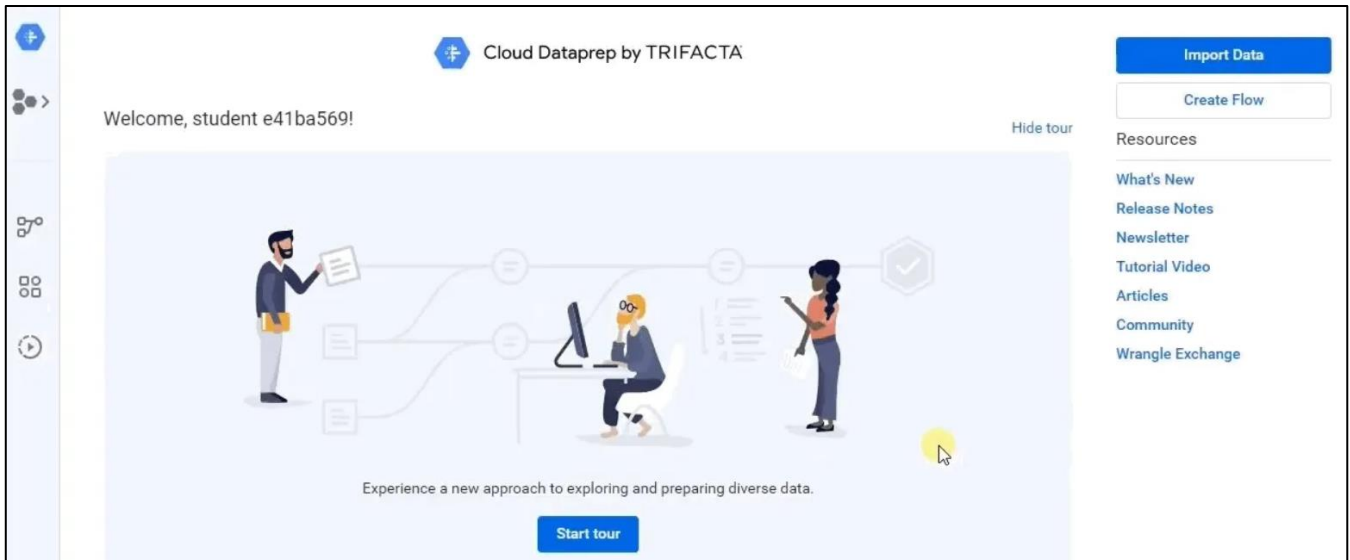
Labels
[+ ADD LABEL](#)

13. Click **CREATE**.

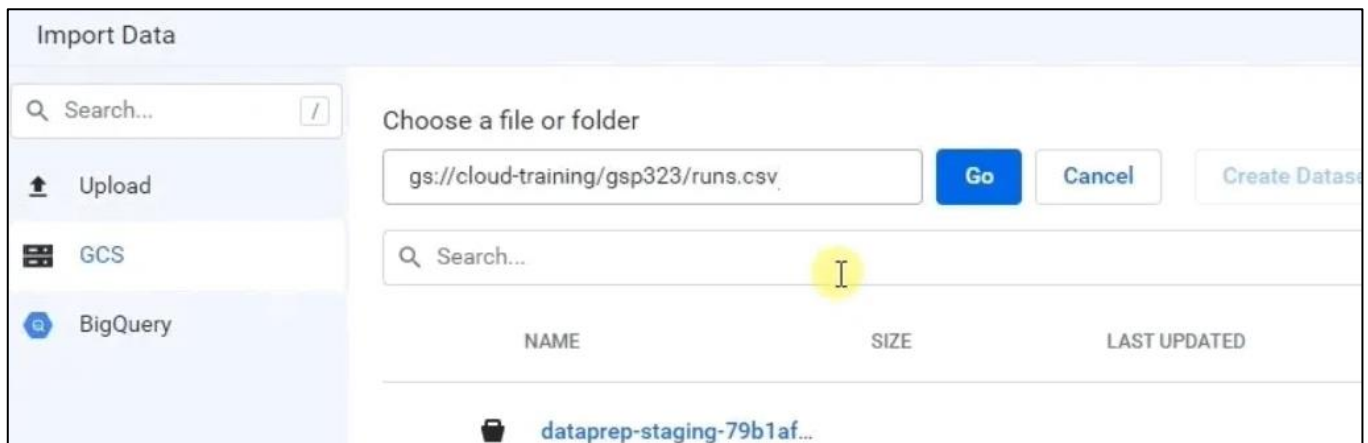
Task 3: Run a simple Dataprep job

Import runs.csv to Dataprep

1. In the Cloud Console, click on **Navigation menu > Dataprep**.
2. After entering the home page of Cloud Dataprep, click the **Import Data** button.



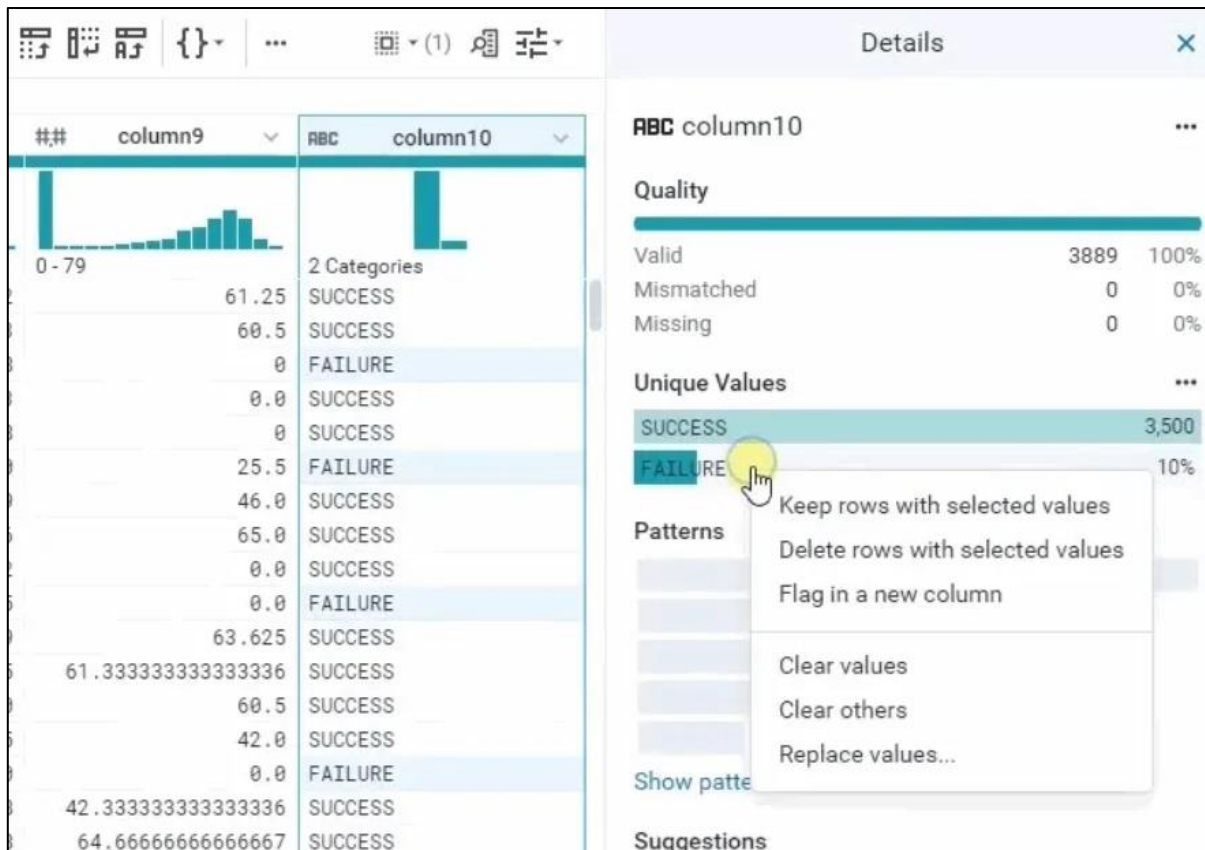
3. In the Import Data page, select **GCS** in the left pane.
4. Click on the pencil icon under Choose a file or folder.
5. Copy `gs://cloud-training/gsp323/runs.csv` to the textbox, and click the **Go** button next to it.



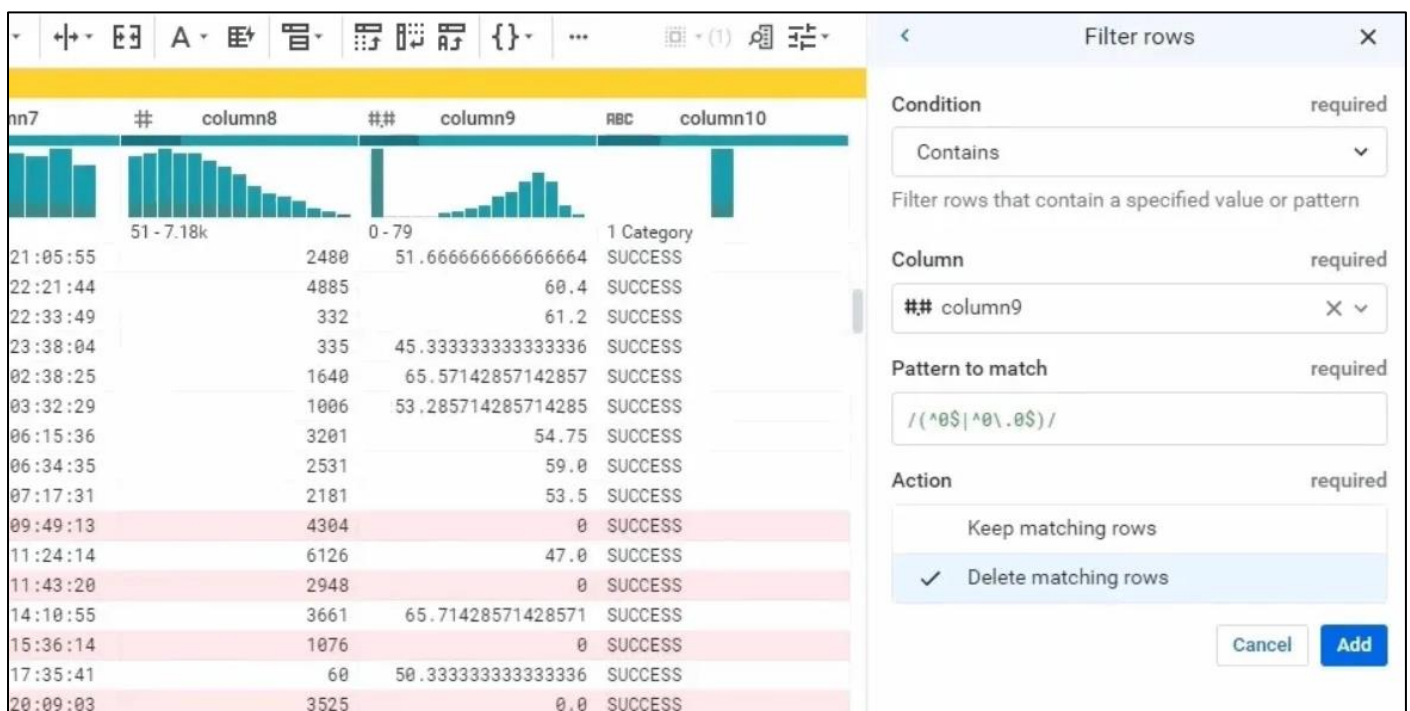
6. After showing the preview of runs.csv in the right pane, click on the **Import & Wrangle** button.

Transform data in Dataprep

1. After launching the Dataprep Transformer, scroll right to the end and select **column10**.
2. In the Details pane, click **FAILURE** under Unique Values to show the context menu.
3. Select **Delete rows with selected values** to Remove all rows with the state of "FAILURE".



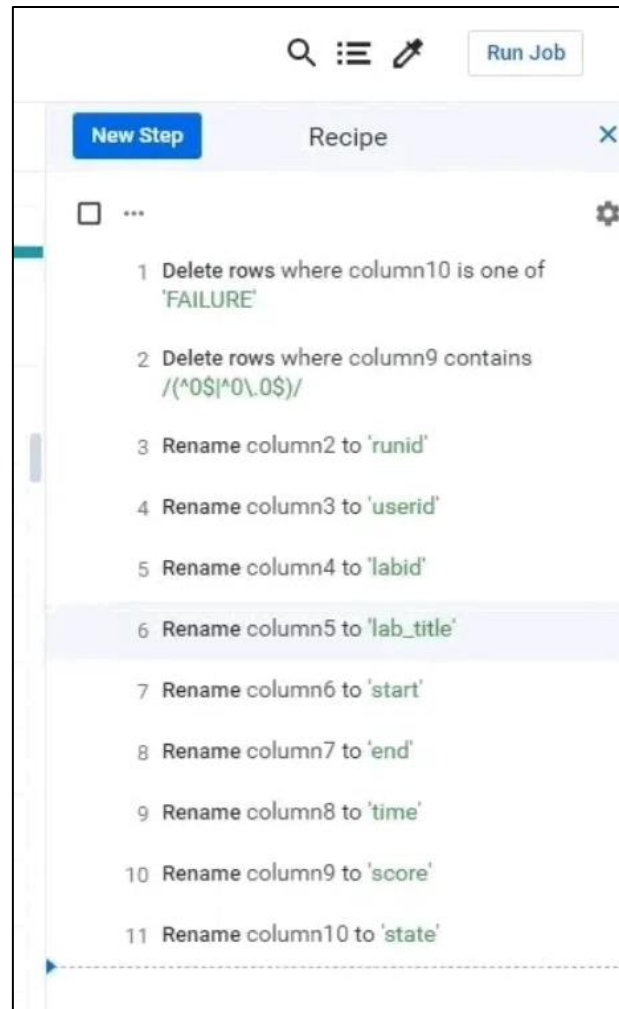
4. Click the downward arrow next to **column9**, choose **Filter rows** > **On column value** > **Contains**.
5. In the Filter rows pane, enter the regex pattern `/(^0$|^0\..0$)/` to "Pattern to match".
6. Select **Delete matching rows** under the Action section, then click the **Add** button.



7. Rename the columns to be:
 - runid
 - userid
 - labid
 - lab_title

- start
- end
- time
- score
- state

8. Confirm the recipe. It should like the screenshot below.



9. Click **Run Job**.

Run Job on Dataflow

Options

☒ Profile results

When enabled, this will generate a profile of your results

Publishing Actions

Add Publishing Action

Actions	Location	Settings
Create-CSV	gs://dataprep-staging-79b1aff1-e834-416d-9c19-ef1fdea335bd/student-00-5a88cc654bc3@qwiklabs.net/jobrun/runs__2.csv	no compression, multiple files, with quotes, with delimiter: ,

Dataflow Execution Settings

Region

us-central1

Zone

Auto Zone

Machine Type

n1-standard-1

Advanced Settings

Task 4: AI

TASK 4 - PART 1 -
CLOUD NATURAL
LANGUAGE:

```
gcloud iam service-accounts create my-natlang-sa \
  --display-name "my natural language service account"

gcloud iam service-accounts keys create ~/key.json \
  --iam-account my-natlang-sa@${GOOGLE_CLOUD_PROJECT}.iam.gserviceaccount.com

export GOOGLE_APPLICATION_CREDENTIALS="/home/$USER/key.json"

gcloud auth activate-service-account my-natlang-
sa@${GOOGLE_CLOUD_PROJECT}.iam.gserviceaccount.com --key-
file=$GOOGLE_APPLICATION_CREDENTIALS

gcloud ml language analyze-entities --content="Old Norse texts portray Odin as
one-eyed and long-bearded, frequently wielding a spear named Gungnir and wearing a
cloak and a broad hat." > result.json

gcloud auth login
(Copy the token from the link provided)

gsutil cp result.json gs://YOUR_PROJECT-marking/task4-cn1.result
```

Use Google Cloud Speech API to analyze the audio file

1. In the Cloud Console, click on **Navigation menu > APIs & Services > Credentials**.
2. In the Credentials page, click on **+ CREATE CREDENTIALS > API key**.
3. Copy the API key to the clipboard, then click **RESTRICT KEY**.
4. Open the Cloud Shell, store the API key as an environment variable by running the following command:

```
export API_KEY=<YOUR-API-KEY>
```

5. **Replace** `<YOUR-API-KEY>` with the copied key value.

```
nano
request.json
```

```
{
  "config": {
    "encoding": "FLAC",
    "languageCode": "en-US"
  },
  "audio": {
    "uri": "gs://cloud-training/gsp323/task4.flac"
  }
}
```

```
curl -s -X POST -H "Content-Type: application/json" --data-binary @request.json \
"https://speech.googleapis.com/v1/speech:recognize?key=${API_KEY}" > result.json
```

```
gsutil cp result.json gs://YOUR_PROJECT-marking/task4-gcs.result
```

```
# TASK 4
- PART 3
- VIDEO
INTELLIGE
NCE:
```

```
gcloud iam service-accounts create quickstart
```

```
gcloud iam service-accounts keys create key.json --iam-account
quickstart@${GOOGLE_CLOUD_PROJECT}.iam.gserviceaccount.com
```

```
gcloud auth activate-service-account --key-file key.json
```

```
export ACCESS_TOKEN=$(gcloud auth print-access-token)
```

```
nano request.json
```

```
{
  "inputUri": "gs://spl/spls/gsp154/video/chicago.mp4",
  "features": [
    "TEXT_DETECTION"
  ]
}
```

```
curl -s -H 'Content-Type: application/json' \
-H "Authorization: Bearer $ACCESS_TOKEN" \
'https://videointelligence.googleapis.com/v1/videos:annotate' \
-d @request.json
```

```
curl -s -H 'Content-Type: application/json' -H "Authorization:
Bearer $ACCESS_TOKEN"
'https://videointelligence.googleapis.com/v1/operations/OPERATION_FR
OM_PREVIOUS_REQUEST' > result1.json

gsutil cp result1.json gs://YOUR_PROJECT-marking/task4-gvi.result
```

Congratulations! You completed this challenge lab.